



CLOUD **FOUNDRY**

# Pivotal Cloud Foundry Developer

## Lab Instructions

Application Deployment using Pivotal Cloud  
Foundry

Version 1.11.a

**Pivotal**

# Copyright Notice

- Copyright © 2017 Pivotal Software, Inc. All rights reserved. This manual and its accompanying materials are protected by U.S. and international copyright and intellectual property laws.
- Pivotal products are covered by one or more patents listed at <http://www.pivotal.io/patents>.
- Pivotal is a registered trademark or trademark of Pivotal Software, Inc. in the United States and/or other jurisdictions. All other marks and names mentioned herein may be trademarks of their respective companies. The training material is provided “as is,” and all express or implied conditions, representations, and warranties, including any implied warranty of merchantability, fitness for a particular purpose or non-infringement, are disclaimed, even if Pivotal Software, Inc., has been advised of the possibility of such claims. This training material is designed to support an instructor-led training course and is intended to be used for reference purposes in conjunction with the instructor-led training course. The training material is not a standalone training tool. Use of the training material for self-study without class attendance is not recommended.
- These materials and the computer programs to which it relates are the property of, and embody trade secrets and confidential information proprietary to, Pivotal Software, Inc., and may not be reproduced, copied, disclosed, transferred, adapted or modified without the express written approval of Pivotal Software, Inc.



---

# Table of Contents

<b>1. Application Autoscaler .....</b>	<b>1</b>
1.1. Preface .....	1
1.2. Setup .....	1
1.3. Provision and Bind an Autoscaler Service Instance .....	1
1.4. Configuring Autoscaling Service Instance .....	2
1.5. Generate Load and Observe the Results .....	3
1.5.1. Generating a Load with JMeter .....	4
1.6. Clean up .....	6
1.7. Questions .....	6

---

# Chapter 1. Application Autoscaler

*Estimated Time: 25 minutes*

## 1.1. Preface

So you can deploy your app, you can scale it. But what if we wish to automate scaling an application up and down during periods of higher and subsequently lower traffic. PCF allows us to automate scaling via the marketplace service named the *App Autoscaler*.

In this exercise we'll introduce a load against our application, and watch it automatically scale up to multiple instances, and then automatically scale back down to a specified minimum. Enjoy.

## 1.2. Setup

1. Ensure the `articulate` application has one application instance running.

```
cf scale articulate -i 1
```

## 1.3. Provision and Bind an Autoscaler Service Instance

1. Read the documentation about [App Autoscaling](#).
2. Review what's in the marketplace.

```
cf marketplace
```

3. Create a autoscaler service instance.

```
cf create-service app-autoscaler standard autoscaler
```

4. Bind the service to `articulate`.

```
cf bind-service articulate autoscaler
```



### Note

You can ignore the "TIP: Use 'cf restage articulate' to ensure your env variable changes take effect" message at this time.

5. Restart the application.

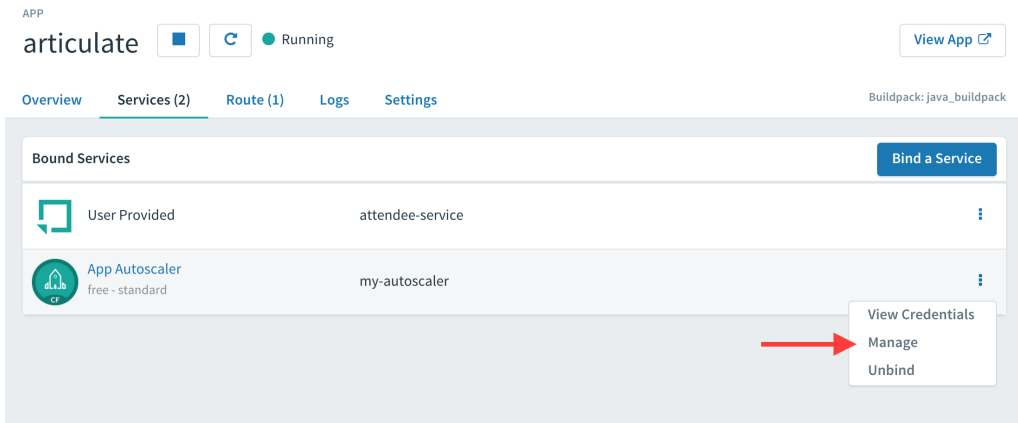
```
cf restart articulate
```

## 1.4. Configuring Autoscaling Service Instance

1. View `articulate` in Apps Manager.

a. Navigate to the `Services` tab.

b. Click the `Manage` link for the autoscaler (this will open a new tab).

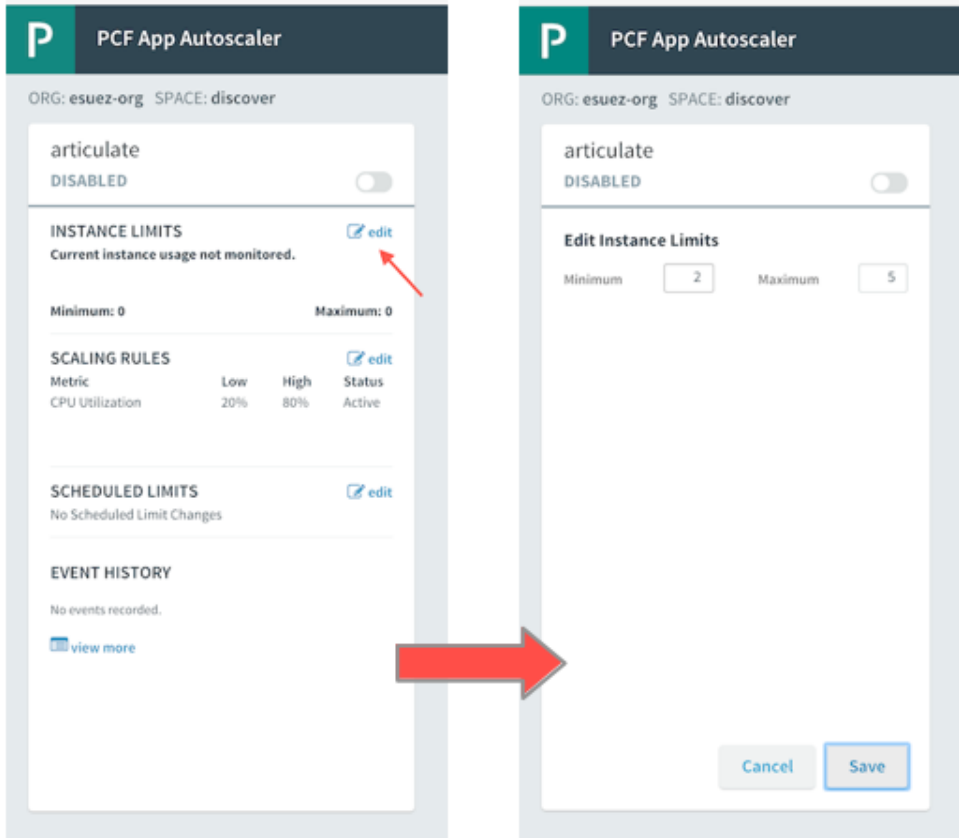


2. Edit instance limits

a. Click 'Edit'

b. Set the minimum number of instances to 2 and the maximum to 5

c. Click 'Save'



3. Return to Apps Manager and observe the number of instances increase from 1 to 2.

## 1.5. Generate Load and Observe the Results

One popular option for generating a load against a web site is by using the Apache load-testing tool [JMeter](#). Instructions for JMeter are outlined below.



## Note

Another option is to use a PWS marketplace service named *Load Impact*. One advantage of using Load Impact is that you don't have to download and install any software to your laptop in order to conduct a load test: it's all done via a web interface. Further, the load test itself is run from the cloud: the test specification can specify what geographic location will serve as the origin of the load test. If you feel like trying Load Impact instead of JMeter, head on over to the PWS Marketplace in your Apps Manager, and create an instance of that service, then follow the instructions for creating a test.

### 1.5.1. Generating a Load with JMeter

1. You should already have `apache-jmeter-XXX.zip` in your lab files, where XXX is the version number. (if you prefer you can download the latest [Apache JMeter](#) for yourself). It will be used to generate load. Review the JMeter [getting started](#) directions.

2. Do *ONE* of the following:

#### Linux and Mac

- a. Add execute permissions to the JMeter script: `chmod +x {{jmeter_basedir}}/bin/jmeter.sh`
- b. Start JMeter by executing `{{jmeter_basedir}}/bin/jmeter.sh`

#### Windows

- a. Start JMeter by double clicking the following file in Windows Explorer:  
`{{jmeter_basedir}}/bin/jmeter.bat`
3. In your lab files you will find a JMeter configuration file: `load-gen.jmx`.
  4. Import the `load-gen.jmx` file with the JMeter GUI (File -> Open).
  5. Expand the content of the JMeter test plan by:
    - a. Selecting the Load Generator test plan on the left pane.
    - b. Going to the menu: Options -> Expand All.
  6. Under HTTP Request, edit the test plan as follows:
    - a. Set the value of the Service Name or IP field to your application's fully-qualified hostname (do not



include the `http://`). For example: `jsmith-articulate.cfapps.io`

b. Set the value of the HTTP Request Path field to `/services`

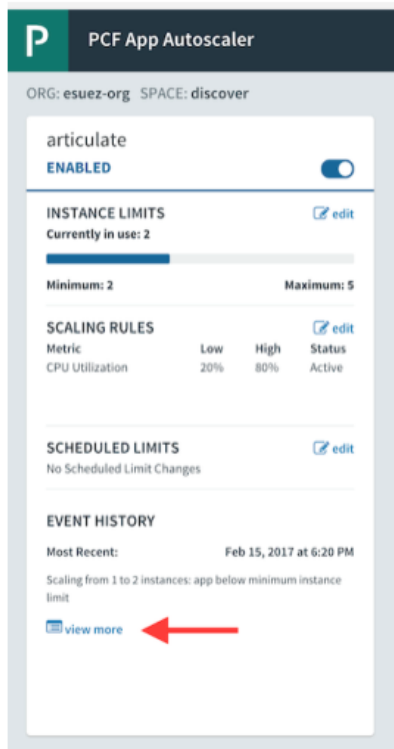
7. Save the test plan: File -> Save.

8. Run the test plan: Run -> Start.

9. Use Apps Manager and the `cf` CLI to observe your service scale up and back down based on load.

10. Stop the test plan: Run -> Stop.

11. Review Autoscale history:



**Figure 1.1. History**



## Note

An even better option for monitoring the state of your application, including your application instances and their CPU usage, is to install the cf cli plugin **cf top**. Head on over to the cloud foundry [plugins catalog](#), locate the plugin named 'top' and install it. It's definitely worth your time.

## 1.6. Clean up

1. Unbind the autoscaler service instance.

```
cf unbind-service articulate autoscaler
```

2. Delete the autoscaler service instance.

```
cf delete-service autoscaler
```

3. Scale articulate back to original settings.

```
cf scale articulate -i 1
```

4. Restart articulate.

```
cf restart articulate
```

## 1.7. Questions

- How do you handle autoscaling today?
- What 12 factor principles are important when it comes to scaling?
- How do you handle scaling at the data layer?